

Mnemorph

An agent harness where trust is earned, proven and written down

An agent can do anything, and nothing proves it did

Power comes from a prompt

What an agent may do lives in text it can read, and often rewrite.

Done is its own word

It says the work is finished. Nothing outside it has checked.

Mistakes leave no trail

Logs live where the agent writes, so they can be changed afterwards.

Mnemorph makes the rules code, not configuration

Every agent is a creature of a race. What each race may do, whom it answers to and how it rises or falls is compiled in.

Rust core

The law is compiled. No prompt or config file can widen it.

Kernel sandbox

Every command runs under a guard, in the soil its creature owns.

Trusted journals

Hash-chained, kept outside the world, out of every body's reach.

A COLONY, WITH ONE JOB FOR EACH CREATURE



Rimuru

Understands you, plans and answers.



Elf

Plans one lane of work.



Orc

Writes the test first, judges the landing.



Slime

The only race that writes code.



Kijin

The only hands that reach outside.



Dark Elf

Attacks the code. Proves every claim.

You stand over all of them as Veldora. Five ascended forms are earned, and can be lost.

Nothing lands because an agent says so

Plan
exactly what would happen

→

Gate
you say yes, bound to that plan

→

Act
in the sandbox

→

Probe
read back from outside

→

Receipt
journaled, undoable

The end gate

Runs the verify lines you pinned. A body cannot edit them.

A cut run

Is probed, never replayed. If it landed, it isn't done twice.

A changed world

Runs nothing. Your yes was for what you saw.

The colony evolves, on trial

Due
stressed, desk full

→

Draft
a model writes traits only

→

Prove
on a copy of the world

→

Trial
10 runs vs doing nothing

→

Keep or roll back
by the rule

Kept only if it wins

At least 2 more tasks won than lost, with no more problems than doing nothing.

Never self-rewriting

The harness writes the file. Race, ground and verify lines are out of reach.

Measured on real memory, recall needed fixing

0.03

Before: F1 at 101 notes

A fixed setting let nearly every note match once memory grew.

0.66

Best fit: F1 at 101 notes

The right setting grows with the square root of memory size.

How we know

Notes written by real bench runs, each task's ask as the query. One model, small tasks: a first calibration.

The board shows the colony, live



The net: a 3D city of races, bodies and Minds

Approvals with the terminal's protections

Memory, ontology, metrics for every world

Loopback only: nothing leaves your machine

Honest numbers: a plain ask now works

10/10

Auto

A plain ask in a full colony, up from 3/10: hidden tests passed.

10/10

Solo

The throne alone, at 60% of auto's cost. Auto has yet to beat it.

4/7

Forced hierarchy

Told to hand everything down: 7 of the 10 tasks ran, 4 passed.

One run per task, one cheap model: a signal, not a proof.

Proven against the real world, not only tests

A live forge run

A branch and a pull request through Kijin, one approval. It found a login bug the stub forge had hidden.

CI on two systems

macOS caught a journal path bug that Linux could not show.

A real audit

Lockfiles checked offline against OSV, in the sandbox, network off.

Where it stands

M0–M13

Built

Core, sandbox, memory, the world, Kijin, Dark Elf, evolution, the board.

1.0

Next

A signed release: a fresh machine running in under two minutes. The hierarchy earns its place task by task.

The SDK and protocol are on npm: @mnemorph/sdk and @mnemorph/protocol